

EE_627_Dylan_Zenner_HW_9

November 21, 2021

```
[1]: !pip3 install pyspark
```

Requirement already satisfied: pyspark in ./opt/anaconda3/lib/python3.8/site-packages (3.2.0)

Requirement already satisfied: py4j==0.10.9.2 in ./opt/anaconda3/lib/python3.8/site-packages (from pyspark) (0.10.9.2)

```
[2]: from pyspark.sql import SparkSession
from pyspark.sql.functions import col
from pyspark.sql.functions import lit
from pyspark.ml.feature import OneHotEncoder, StringIndexer, VectorAssembler
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
```

```
[3]: spark = SparkSession.builder.appName('HW 9').getOrCreate()
```

21/11/21 18:02:00 WARN Utils: Your hostname, dylans-MacBook-Pro.local resolves to a loopback address: 127.0.0.1; using 192.168.1.125 instead (on interface en0)
21/11/21 18:02:00 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address

Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".

To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).

21/11/21 18:02:00 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

```
[4]: spark
```

```
[4]: <pyspark.sql.session.SparkSession at 0x7ff1633c84c0>
```

```
[5]: data_df = pd.read_csv("test2_new.txt", sep='|', names=['userID', 'trackID', 'predictions'])
data_df.head()
```

```
[5]:   userID  trackID  predictions
0  200031    30877             1
```

1	200031	8244	1
2	200031	130183	0
3	200031	198762	0
4	200031	34503	1

```
[6]: data2_df = pd.read_csv("output1.txt", sep='|', names=['userID', 'trackID', 'album_score', 'artist_score'])
data2_df.head()
```

```
[6]:   userID  trackID  album_score  artist_score
0  199810   208019           0.0           0.0
1  199810    74139           0.0           0.0
2  199810    9903           0.0           0.0
3  199810   242681           0.0           0.0
4  199810    18515           0.0          70.0
```

```
[7]: data2_df.columns[data2_df.isnull().any()] # check for null values
```

```
[7]: Index([], dtype='object')
```

```
[8]: scores_df = data_df.merge(data2_df, how='inner', on=['userID', 'trackID']) # inner join the dataframes
scores_df.head()
```

```
[8]:   userID  trackID  predictions  album_score  artist_score
0  200031   30877           1           90.0           50.0
1  200031    8244           1           90.0           0.0
2  200031   130183          0           0.0           0.0
3  200031   198762          0           0.0           0.0
4  200031    34503           1           90.0           50.0
```

```
[9]: scores_df.columns[scores_df.isnull().any()] # check for null values
```

```
[9]: Index([], dtype='object')
```

```
[10]: scores_df.to_csv('ratings.csv', index=False)
scores_df.count()
```

```
[10]: userID          6000
trackID          6000
predictions      6000
album_score      6000
artist_score     6000
dtype: int64
```

```
[11]: scores_df = spark.read.csv('ratings.csv', header=True, inferSchema=True)
scores_df.printSchema()
```

```

root
|-- userID: integer (nullable = true)
|-- trackID: integer (nullable = true)
|-- predictions: integer (nullable = true)
|-- album_score: double (nullable = true)
|-- artist_score: double (nullable = true)

```

```

[12]: numericCols = ['album_score', 'artist_score']
      stages = []
      assemblerInputs = numericCols
      assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
      stages += [assembler]

```

```

[13]: label_stringIdx = StringIndexer(inputCol = 'predictions', outputCol = 'label')
      stages += [label_stringIdx]

```

```

[14]: from pyspark.ml import Pipeline
      pipeline = Pipeline(stages = stages)
      pipelineModel = pipeline.fit(scores_df)
      df = pipelineModel.transform(scores_df)
      cols = scores_df.columns
      selectedCols = ['label', 'features'] + cols
      df = df.select(selectedCols)
      df.printSchema()

```

```

root
|-- label: double (nullable = false)
|-- features: vector (nullable = true)
|-- userID: integer (nullable = true)
|-- trackID: integer (nullable = true)
|-- predictions: integer (nullable = true)
|-- album_score: double (nullable = true)
|-- artist_score: double (nullable = true)

```

```

[15]: pd.DataFrame(df.take(5), columns=df.columns).transpose()

```

```

[15]:
      0      1      2      3      4
label      1.0      1.0      0.0      0.0      1.0
features  [90.0, 50.0]  [90.0, 0.0]  (0.0, 0.0)  (0.0, 0.0)  [90.0, 50.0]
userID      200031      200031      200031      200031      200031
trackID      30877      8244      130183      198762      34503
predictions      1      1      0      0      1
album_score      90.0      90.0      0.0      0.0      90.0

```

artist_score	50.0	0.0	0.0	0.0	50.0
--------------	------	-----	-----	-----	------

```
[16]: # below is the typical random split
# of the train and test data sets
# HOWEVER, our testing users have 6 tracks for each
# We cannot use random split here
train, test = df.randomSplit([0.7, 0.3], seed = 2018)
print("Training Dataset Count: " + str(train.count()))
print("Test Dataset Count: " + str(test.count()))
```

Training Dataset Count: 4260

Test Dataset Count: 1740

```
[17]: predictions_df = spark.read.csv('output1.txt', sep='|', inferSchema=True)
```

```
[18]: predictions_df.count()
```

```
[18]: 120000
```

```
[19]: predictions_df
```

```
[19]: DataFrame[_c0: int, _c1: int, _c2: double, _c3: double]
```

```
[20]: predictions_df = predictions_df.withColumnRenamed("_c0", "userID").
      ↳withColumnRenamed("_c1", "trackID").withColumnRenamed("_c2", "album_score").
      ↳withColumnRenamed("_c3", "artist_score")
```

```
[21]: predictions_columns = predictions_df.columns
      predictions_columns
```

```
[21]: ['userID', 'trackID', 'album_score', 'artist_score']
```

```
[22]: predictions_df = predictions_df.withColumn('prediction', lit('0'))
```

```
[23]: pd.DataFrame(predictions_df.take(10), columns=predictions_df.columns).
      ↳transpose()
```

```
[23]:
```

	0	1	2	3	4	5	6	7 \
userID	199810	199810	199810	199810	199810	199810	199812	199812
trackID	208019	74139	9903	242681	18515	105760	276940	142408
album_score	0.0	0.0	0.0	0.0	0.0	0.0	0.0	100.0
artist_score	0.0	0.0	0.0	0.0	70.0	90.0	0.0	100.0
prediction	0	0	0	0	0	0	0	0

	8	9
userID	199812	199812
trackID	130023	29189
album_score	100.0	0.0

```

artist_score    100.0    0.0
prediction       0        0

```

```
[24]: predictions_df.printSchema()
```

```

root
|-- userID: integer (nullable = true)
|-- trackID: integer (nullable = true)
|-- album_score: double (nullable = true)
|-- artist_score: double (nullable = true)
|-- prediction: string (nullable = false)

```

```
[25]: numericCols = ['album_score', 'artist_score']
      stages = []
      assemblerInputs = numericCols
      assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
      stages += [assembler]
```

```
[26]: label_stringIdx = StringIndexer(inputCol = 'prediction', outputCol = 'label')
      stages += [label_stringIdx]
```

```
[27]: predictions_pipeline = Pipeline(stages=stages)
      predictions_pipeline_model = predictions_pipeline.fit(predictions_df)
      predictions_df = predictions_pipeline_model.transform(predictions_df)
```

```
[28]: selected_columns = ['label', 'features'] + predictions_columns
      predictions_df = predictions_df.select(selected_columns)
      predictions_df.printSchema()
```

```

root
|-- label: double (nullable = false)
|-- features: vector (nullable = true)
|-- userID: integer (nullable = true)
|-- trackID: integer (nullable = true)
|-- album_score: double (nullable = true)
|-- artist_score: double (nullable = true)

```

```
[29]: pd.DataFrame(predictions_df.take(10), columns=predictions_df.columns).
      ↪transpose()
```

```
[29]:
```

	0	1	2	3	4 \
label	0.0	0.0	0.0	0.0	0.0
features	(0.0, 0.0)	(0.0, 0.0)	(0.0, 0.0)	(0.0, 0.0)	[0.0, 70.0]
userID	199810	199810	199810	199810	199810
trackID	208019	74139	9903	242681	18515
album_score	0.0	0.0	0.0	0.0	0.0

artist_score	0.0	0.0	0.0	0.0	70.0
--------------	-----	-----	-----	-----	------

	5	6	7	8	\
label	0.0	0.0	0.0	0.0	
features	[0.0, 90.0]	(0.0, 0.0)	[100.0, 100.0]	[100.0, 100.0]	
userID	199810	199812	199812	199812	
trackID	105760	276940	142408	130023	
album_score	0.0	0.0	100.0	100.0	
artist_score	90.0	0.0	100.0	100.0	


	9
label	0.0
features	(0.0, 0.0)
userID	199812
trackID	29189
album_score	0.0
artist_score	0.0

1 Logistic regression model

```
[30]: from pyspark.ml.classification import LogisticRegression
```


```
[31]: logistic_regression = LogisticRegression(featuresCol = 'features', labelCol = 'label',)
logistic_regression_model = logistic_regression.fit(train)
```

```
[32]: beta = np.sort(logistic_regression_model.coefficients)
plt.plot(beta)
plt.ylabel('Beta Coefficients')
plt.show()
```



```
[33]: trainingSummary = logistic_regression_model.summary
      roc = trainingSummary.roc.toPandas()
      plt.plot(roc['FPR'],roc['TPR'])
      plt.ylabel('False Positive Rate')
      plt.xlabel('True Positive Rate')
      plt.title('ROC Curve')
      plt.show()
      print('Training set areaUnderROC: ' + str(trainingSummary.areaUnderROC))

/Users/dylan/opt/anaconda3/lib/python3.8/site-
packages/pyspark/sql/context.py:125: FutureWarning: Deprecated in 3.0.0. Use
SparkSession.builder.getOrCreate() instead.
  warnings.warn(
```



Training set areaUnderROC: 0.8737133978552075

```
[34]: predictions = logistic_regression_model.transform(test)
      predictions.select('userID', 'trackID', 'label', 'probability',
                        'rawPrediction', 'prediction' ).show(12)
```

userID	trackID	label	probability	rawPrediction	prediction
200031	130183	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200065	179571	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200070	124239	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200070	271459	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200085	134106	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200099	41892	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200106	152491	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200124	284066	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200143	131171	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200143	187136	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200160	231680	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200166	193878	0.0	[0.80218276453714...	[1.39999291666030...	0.0

only showing top 12 rows


```
[35]: # col("userID").asc() sort the user ascending
# col("probability").desc() sort the probability descending (from large to
      ↪small)
from pyspark.sql.functions import col
sort_predictions = predictions.select('userID', 'trackID',
                                     'label', 'probability',
                                     'rawPrediction', 'prediction'
                                     ).sort(col("userID").asc(),
      ↪col("probability").desc())
sort_predictions.show(6)
```

userID	trackID	label	probability	rawPrediction	prediction
200031	130183	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200031	8244	1.0	[0.06488875729218...	[-2.6679911209111...	1.0
200031	30877	1.0	[0.01343471003408...	[-4.2963878477543...	1.0
200055	56695	1.0	[0.00368752576332...	[-5.5991052292289...	1.0
200065	179571	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200065	119451	1.0	[0.17782970128751...	[-1.5311211916575...	1.0

only showing top 6 rows

```
[36]: # probability column is where you can decide
# the 3 tracks with highest "probability" for "1" and
# the 3 tracks with lowest "probability" for "0"
# tracks "198762", "34503", "130183" with "1"
# tracks "30877", "8244", "227283" with "0"
sort_predictions.show(6)
```

userID	trackID	label	probability	rawPrediction	prediction
200031	130183	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200031	8244	1.0	[0.06488875729218...	[-2.6679911209111...	1.0
200031	30877	1.0	[0.01343471003408...	[-4.2963878477543...	1.0
200055	56695	1.0	[0.00368752576332...	[-5.5991052292289...	1.0
200065	179571	0.0	[0.80218276453714...	[1.39999291666030...	0.0
200065	119451	1.0	[0.17782970128751...	[-1.5311211916575...	1.0

only showing top 6 rows

```
[37]: logistic_regression_model_predictions = logistic_regression_model.
      ↪transform(predictions_df)
logistic_regression_model_predictions.select('userID', 'trackID',
      ↪'probability', 'rawPrediction', 'prediction').show(12)
```

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80218276453714...	[1.39999291666030...	0.0
199810	74139	[0.80218276453714...	[1.39999291666030...	0.0
199810	9903	[0.80218276453714...	[1.39999291666030...	0.0
199810	242681	[0.80218276453714...	[1.39999291666030...	0.0
199810	18515	[0.29322699696820...	[-0.8797625009202...	1.0
199810	105760	[0.17782970128751...	[-1.5311211916575...	1.0
199812	276940	[0.80218276453714...	[1.39999291666030...	0.0
199812	142408	[0.00169769820591...	[-6.3767828009944...	1.0
199812	130023	[0.00169769820591...	[-6.3767828009944...	1.0
199812	29189	[0.80218276453714...	[1.39999291666030...	0.0
199812	223706	[0.13507641145138...	[-1.8568005370262...	1.0
199812	211361	[0.80218276453714...	[1.39999291666030...	0.0

only showing top 12 rows

```
[38]: sort_logistic_regression_model_predictions =
↳ logistic_regression_model_predictions.select('userID', 'trackID',
↳ 'probability', 'rawPrediction', 'prediction').sort(col('userID').asc(),
↳ col('probability').desc())
sort_logistic_regression_model_predictions.show(6)
```

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80218276453714...	[1.39999291666030...	0.0
199810	242681	[0.80218276453714...	[1.39999291666030...	0.0
199810	74139	[0.80218276453714...	[1.39999291666030...	0.0
199810	9903	[0.80218276453714...	[1.39999291666030...	0.0
199810	18515	[0.29322699696820...	[-0.8797625009202...	1.0
199810	105760	[0.17782970128751...	[-1.5311211916575...	1.0

only showing top 6 rows

```
[39]: sort_logistic_regression_model_predictions_df =
↳ sort_logistic_regression_model_predictions.toPandas()
sort_logistic_regression_model_predictions_df.
↳ columns[sort_logistic_regression_model_predictions_df.isnull().any()]
```

```
[39]: Index([], dtype='object')
```

```
[40]: sort_logistic_regression_model_predictions_df.head()
```

```
[40]:   userID  trackID      probability \
0  199810   208019  [0.8021827645371431, 0.1978172354628569]
1  199810    74139  [0.8021827645371431, 0.1978172354628569]
2  199810     9903  [0.8021827645371431, 0.1978172354628569]
3  199810   242681  [0.8021827645371431, 0.1978172354628569]
4  199810    18515  [0.2932269969682032, 0.7067730030317968]

      rawPrediction  prediction
0  [1.3999929166603091, -1.3999929166603091]      0.0
1  [1.3999929166603091, -1.3999929166603091]      0.0
2  [1.3999929166603091, -1.3999929166603091]      0.0
3  [1.3999929166603091, -1.3999929166603091]      0.0
4  [-0.8797625009202483, 0.8797625009202483]      1.0
```

```
[41]: columns = ['userID', 'trackID']
sort_logistic_regression_model_predictions_df.
      ↳to_csv('logistic_regression_predictions.csv', index=False, header=None,
      ↳columns=columns)
```

```
[42]: logistic_regression_predictions_file = open('logistic_regression_predictions.
      ↳csv')
final_logistic_regression_predictions_file =
      ↳open('logistic_regression_submission.csv', 'w')
```

```
[43]: final_logistic_regression_predictions_file.write('TrackID,Predictor\n')
```

```
[43]: 18
```

```
[44]: last_user_id = -1
track_id_out_vec = [0] * 6
```

```
[45]: for line in logistic_regression_predictions_file:
      arr_out = line.strip().split(',')
      user_id_out = arr_out[0]
      track_id_out = arr_out[1]

      if user_id_out != last_user_id:
          i = 0

      track_id_out_vec[i] = track_id_out

      i = i + 1
      last_user_id = user_id_out

      if i == 6:
```

```

        predictions = np.ones(shape=(6))
        for index in range(0, 3):
            predictions[index] = 0

        for ii in range(0, 6):
            out_str = str(user_id_out) + '_' + str(track_id_out_vec[ii]) + ',' +
↪+ str(int(predictions[ii]))
            final_logistic_regression_predictions_file.write(out_str + '\n')

```

```

[46]: logistic_regression_predictions_file.close()
      final_logistic_regression_predictions_file.close()

```

2 Decision Tree Classifier

```

[47]: from pyspark.ml.classification import DecisionTreeClassifier

```

```

[48]: decision_tree_classifier = DecisionTreeClassifier(featuresCol='features',
↪labelCol='label', maxDepth=3)
      decision_tree_classifier_model = decision_tree_classifier.fit(train)

```

```

[49]: predictions_decision_tree_classifier = decision_tree_classifier_model.
↪transform(test)

```

```

[50]: evaluator = MulticlassClassificationEvaluator(labelCol='label',
↪predictionCol='prediction', metricName='accuracy')
      accuracy = evaluator.evaluate(predictions_decision_tree_classifier)
      print('Test Error: {}'.format((1.0 - accuracy) * 100))

```

Test Error: 14.482758620689651

```

[51]: sort_predictions_decision_tree_classifier =
↪predictions_decision_tree_classifier.select('userID', 'trackID', 'label',
↪'probability', 'rawPrediction', 'prediction').sort(col('userID').asc(),
↪col('probability').desc())
      sort_predictions_decision_tree_classifier.show(6)

```

```

+-----+-----+-----+-----+-----+-----+
|userID|trackID|label|      probability| rawPrediction|prediction|
+-----+-----+-----+-----+-----+-----+
|200031| 130183| 0.0|[0.80858085808580...|[1960.0,464.0]|      0.0|
|200031|  30877| 1.0|[0.08880090497737...|[157.0,1611.0]|      1.0|
|200031|   8244| 1.0|[0.01470588235294...|  [1.0,67.0]|      1.0|
|200055|  56695| 1.0|[0.08880090497737...|[157.0,1611.0]|      1.0|
|200065| 179571| 0.0|[0.80858085808580...|[1960.0,464.0]|      0.0|
|200065| 119451| 1.0|[0.08880090497737...|[157.0,1611.0]|      1.0|
+-----+-----+-----+-----+-----+-----+
only showing top 6 rows

```

```
[52]: predictions_decision_tree_classifier = decision_tree_classifier_model.
      ↪transform(predictions_df)
      predictions_decision_tree_classifier.select('userID', 'trackID', 'probability',
      ↪'rawPrediction', 'prediction').show(12)
```

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80858085808580...	[1960.0,464.0]	0.0
199810	74139	[0.80858085808580...	[1960.0,464.0]	0.0
199810	9903	[0.80858085808580...	[1960.0,464.0]	0.0
199810	242681	[0.80858085808580...	[1960.0,464.0]	0.0
199810	18515	[0.08880090497737...	[157.0,1611.0]	1.0
199810	105760	[0.08880090497737...	[157.0,1611.0]	1.0
199812	276940	[0.80858085808580...	[1960.0,464.0]	0.0
199812	142408	[0.08880090497737...	[157.0,1611.0]	1.0
199812	130023	[0.08880090497737...	[157.0,1611.0]	1.0
199812	29189	[0.80858085808580...	[1960.0,464.0]	0.0
199812	223706	[0.08880090497737...	[157.0,1611.0]	1.0
199812	211361	[0.80858085808580...	[1960.0,464.0]	0.0

only showing top 12 rows

```
[53]: sort_predictions_decision_tree_classifier =
      ↪predictions_decision_tree_classifier.select('userID', 'trackID',
      ↪'probability', 'rawPrediction', 'prediction').sort(col('userID').asc(),
      ↪col('probability').desc())
      sort_predictions_decision_tree_classifier.show(6)
```

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80858085808580...	[1960.0,464.0]	0.0
199810	242681	[0.80858085808580...	[1960.0,464.0]	0.0
199810	74139	[0.80858085808580...	[1960.0,464.0]	0.0
199810	9903	[0.80858085808580...	[1960.0,464.0]	0.0
199810	18515	[0.08880090497737...	[157.0,1611.0]	1.0
199810	105760	[0.08880090497737...	[157.0,1611.0]	1.0

only showing top 6 rows

```
[54]: sort_predictions_decision_tree_classifier_df =
      ↪sort_predictions_decision_tree_classifier.toPandas()
      sort_predictions_decision_tree_classifier_df.
      ↪columns[sort_predictions_decision_tree_classifier_df.isnull().any()]
```

```
[54]: Index([], dtype='object')
```

```
[55]: sort_predictions_decision_tree_classifier_df.head()
```

```
[55]:  userID  trackID  probability \
0  199810   208019  [0.8085808580858086, 0.19141914191419143]
1  199810    74139  [0.8085808580858086, 0.19141914191419143]
2  199810    9903  [0.8085808580858086, 0.19141914191419143]
3  199810   242681  [0.8085808580858086, 0.19141914191419143]
4  199810    18515  [0.08880090497737557, 0.9111990950226244]

      rawPrediction  prediction
0  [1960.0, 464.0]          0.0
1  [1960.0, 464.0]          0.0
2  [1960.0, 464.0]          0.0
3  [1960.0, 464.0]          0.0
4  [157.0, 1611.0]          1.0
```

```
[56]: columns = ['userID', 'trackID']
sort_predictions_decision_tree_classifier_df.
      ↳to_csv('decision_tree_classifier_predictions.csv', index=False, header=None,
      ↳columns=columns)
```

```
[57]: decision_tree_classifier_file = open('decision_tree_classifier_predictions.csv')
final_decision_tree_classifier_file = open('decision_tree_classifier_submission.
      ↳csv', 'w')
```

```
[58]: final_decision_tree_classifier_file.write('TrackID,Predictor\n')
```

```
[58]: 18
```

```
[59]: last_user_id = -1
track_id_out_vec = [0] * 6
```

```
[60]: for line in decision_tree_classifier_file:
      arr_out = line.strip().split(',')
      user_id_out = arr_out[0]
      track_id_out = arr_out[1]

      if user_id_out != last_user_id:
          i = 0

      track_id_out_vec[i] = track_id_out

      i = i + 1
      last_user_id = user_id_out
      if i == 6:
          predictions = np.ones(shape=(6))
```

```

    for index in range(0, 3):
        predictions[index] = 0

    for ii in range(0, 6):
        out_str = str(user_id_out) + '_' + str(track_id_out_vec[ii]) + ',' +
        ↪+ str(int(predictions[ii]))
        final_decision_tree_classifier_file.write(out_str + '\n')

```

```

[61]: decision_tree_classifier_file.close()
      final_decision_tree_classifier_file.close()

```

3 Random Forest Classifier

```

[62]: from pyspark.ml.classification import RandomForestClassifier

```

```

[63]: random_forest = RandomForestClassifier(featuresCol='features', labelCol='label')
      random_forest_model = random_forest.fit(train)

```

```

[64]: random_forest_predictions = random_forest_model.transform(test)

```

```

[65]: evaluator = MulticlassClassificationEvaluator(labelCol='label',
        ↪predictionCol='prediction', metricName='accuracy')
      accuracy = evaluator.evaluate(random_forest_predictions)
      print('Test Error: {}'.format((1.0 - accuracy) * 100))

```

Test Error: 14.482758620689651

```

[66]: sort_random_forest_predictions = random_forest_predictions.select('userID',
        ↪'trackID', 'label', 'probability', 'rawPrediction', 'prediction').
        ↪sort(col('userID').asc(), col('probability').desc())
      sort_random_forest_predictions.show(6)

```

```

+-----+-----+-----+-----+-----+-----+
|userID|trackID|label|      probability|      rawPrediction|prediction|
+-----+-----+-----+-----+-----+-----+
|200031| 130183| 0.0|[0.80878408351365...|[16.1756816702730...|    0.0|
|200031|  30877| 1.0|[0.01263507637292...|[0.25270152745857...|    1.0|
|200031|   8244| 1.0|[8.06451612903225...|[0.01612903225806...|    1.0|
|200055|  56695| 1.0|[0.01263507637292...|[0.25270152745857...|    1.0|
|200065| 179571| 0.0|[0.80878408351365...|[16.1756816702730...|    0.0|
|200065| 119451| 1.0|[0.18553913276044...|[3.71078265520889...|    1.0|
+-----+-----+-----+-----+-----+-----+
only showing top 6 rows

```

```

[67]: random_forest_predictions = random_forest_model.transform(predictions_df)

```

```
random_forest_predictions.select('userID', 'trackID', 'probability',
↳ 'rawPrediction', 'prediction').show(12)
```

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80878408351365...	[16.1756816702730...	0.0
199810	74139	[0.80878408351365...	[16.1756816702730...	0.0
199810	9903	[0.80878408351365...	[16.1756816702730...	0.0
199810	242681	[0.80878408351365...	[16.1756816702730...	0.0
199810	18515	[0.19467261212938...	[3.89345224258761...	1.0
199810	105760	[0.18553913276044...	[3.71078265520889...	1.0
199812	276940	[0.80878408351365...	[16.1756816702730...	0.0
199812	142408	[0.01263507637292...	[0.25270152745857...	1.0
199812	130023	[0.01263507637292...	[0.25270152745857...	1.0
199812	29189	[0.80878408351365...	[16.1756816702730...	0.0
199812	223706	[0.18553913276044...	[3.71078265520889...	1.0
199812	211361	[0.80878408351365...	[16.1756816702730...	0.0

only showing top 12 rows

```
[68]: sort_random_forest_predictions = random_forest_predictions.select('userID',
↳ 'trackID', 'probability', 'rawPrediction', 'prediction').sort(col('userID').
↳ asc(), col('probability').desc())
sort_random_forest_predictions.show(6)
```

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80878408351365...	[16.1756816702730...	0.0
199810	242681	[0.80878408351365...	[16.1756816702730...	0.0
199810	74139	[0.80878408351365...	[16.1756816702730...	0.0
199810	9903	[0.80878408351365...	[16.1756816702730...	0.0
199810	18515	[0.19467261212938...	[3.89345224258761...	1.0
199810	105760	[0.18553913276044...	[3.71078265520889...	1.0

only showing top 6 rows

```
[69]: sort_random_forest_predictions_df = sort_random_forest_predictions.toPandas()
sort_random_forest_predictions_df.columns[sort_random_forest_predictions_df.
↳ isnull().any()]
```

```
[69]: Index([], dtype='object')
```



```
[70]: sort_random_forest_predictions_df.head()
```

```
[70]:  userID  trackID  probability \
0  199810   208019  [0.8087840835136524, 0.19121591648634775]
1  199810    74139  [0.8087840835136524, 0.19121591648634775]
2  199810     9903  [0.8087840835136524, 0.19121591648634775]
3  199810   242681  [0.8087840835136524, 0.19121591648634775]
4  199810    18515  [0.1946726121293807, 0.8053273878706193]

      rawPrediction  prediction
0  [16.175681670273047, 3.824318329726955]      0.0
1  [16.175681670273047, 3.824318329726955]      0.0
2  [16.175681670273047, 3.824318329726955]      0.0
3  [16.175681670273047, 3.824318329726955]      0.0
4  [3.893452242587614, 16.106547757412386]      1.0
```

```
[71]: columns= ['userID', 'trackID']
sort_random_forest_predictions_df.to_csv('random_forest_classifier_predictions.
→csv', index=False, header=None, columns=columns)
```

```
[72]: random_forest_predictions_file = open('random_forest_classifier_predictions.
→csv')
final_random_forest_predictions_file =
→open('random_forest_classifier_submission.csv', 'w')
```

```
[73]: final_random_forest_predictions_file.write('TrackID,Predictor\n')
```

```
[73]: 18
```

```
[74]: last_user_id = -1
track_id_out_vec = [0] * 6
```

```
[75]: for line in random_forest_predictions_file:
    arr_out = line.strip().split(',')
    user_id_out = arr_out[0]
    track_id_out = arr_out[1]

    if user_id_out != last_user_id:
        i = 0

    track_id_out_vec[i] = track_id_out

    i = i + 1
    last_user_id = user_id_out

    if i == 6:
        predictions = np.ones(shape=(6))
```

```

    for index in range(0, 3):
        predictions[index] = 0

    for ii in range(0, 6):
        out_str = str(user_id_out) + '_' + str(track_id_out_vec[ii]) + ','
        ↪+ str(int(predictions[ii]))
        final_random_forest_predictions_file.write(out_str + '\n')

```

```

[76]: random_forest_predictions_file.close()
      final_random_forest_predictions_file.close()

```

4 Gradient Boosted Tree Classifier

```

[77]: from pyspark.ml.classification import GBTClassifier

```

```

[78]: gradient_boosted_tree_classifier = GBTClassifier(maxIter=100)
      gradient_boosted_tree_classifier_model = gradient_boosted_tree_classifier.
        ↪fit(train)

```

```

[79]: gradient_boosted_tree_classifier_predictions =
        ↪gradient_boosted_tree_classifier_model.transform(test)

```

```

[80]: evaluator = MulticlassClassificationEvaluator(labelCol='label',
        ↪predictionCol='prediction', metricName='accuracy')
      accuracy = evaluator.evaluate(gradient_boosted_tree_classifier_predictions)
      print('Test Error: {}'.format((1.0 - accuracy) * 100))

```

Test Error: 14.482758620689651

21/11/21 18:03:08 WARN InstanceBuilder\$NativeBLAS: Failed to load implementation from:dev.ludovic.netlib.blas.JNIBLAS

21/11/21 18:03:08 WARN InstanceBuilder\$NativeBLAS: Failed to load implementation from:dev.ludovic.netlib.blas.ForeignLinkerBLAS

```

[81]: sort_gradient_boosted_tree_classifier_predictions =
        ↪gradient_boosted_tree_classifier_predictions.select('userID', 'trackID',
        ↪'label', 'probability', 'rawPrediction', 'prediction').sort(col('userID').
        ↪asc(), col('probability').desc())
      sort_gradient_boosted_tree_classifier_predictions.show(6)

```

userID	trackID	label	probability	rawPrediction	prediction
200031	130183	0.0	[0.80853359039673...]	[0.72025491912707...]	0.0
200031	30877	1.0	[0.01165681655269...]	[-2.2200694341913...]	1.0
200031	8244	1.0	[0.01165681655269...]	[-2.2200694341913...]	1.0
200055	56695	1.0	[0.01632026531261...]	[-2.0494463824457...]	1.0
200065	179571	0.0	[0.80853359039673...]	[0.72025491912707...]	0.0

200065	119451	1.0	[0.19777096337535...	[-0.7001422692047...	1.0
--------	--------	-----	----------------------	----------------------	-----

only showing top 6 rows

```
[82]: gradient_boosted_tree_classifier_predictions =
      ↳ gradient_boosted_tree_classifier_model.transform(predictions_df)
      gradient_boosted_tree_classifier_predictions.select('userID', 'trackID',
      ↳ 'probability', 'rawPrediction', 'prediction').show(12)
```

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80853359039673...	[0.72025491912707...	0.0
199810	74139	[0.80853359039673...	[0.72025491912707...	0.0
199810	9903	[0.80853359039673...	[0.72025491912707...	0.0
199810	242681	[0.80853359039673...	[0.72025491912707...	0.0
199810	18515	[0.24633418901301...	[-0.5591299711568...	1.0
199810	105760	[0.19777096337535...	[-0.7001422692047...	1.0
199812	276940	[0.80853359039673...	[0.72025491912707...	0.0
199812	142408	[0.01161609118736...	[-2.2218399428711...	1.0
199812	130023	[0.01161609118736...	[-2.2218399428711...	1.0
199812	29189	[0.80853359039673...	[0.72025491912707...	0.0
199812	223706	[0.12602709898363...	[-0.9682762071735...	1.0
199812	211361	[0.80853359039673...	[0.72025491912707...	0.0

only showing top 12 rows

```
[83]: sort_gradient_boosted_tree_classifier_predictions =
      ↳ gradient_boosted_tree_classifier_predictions.select('userID', 'trackID',
      ↳ 'probability', 'rawPrediction', 'prediction').sort(col('userID').asc(),
      ↳ col('probability').desc())
      sort_gradient_boosted_tree_classifier_predictions.show(6)
```

[Stage 1106:> (0 + 1) / 1]

userID	trackID	probability	rawPrediction	prediction
199810	208019	[0.80853359039673...	[0.72025491912707...	0.0
199810	242681	[0.80853359039673...	[0.72025491912707...	0.0
199810	74139	[0.80853359039673...	[0.72025491912707...	0.0
199810	9903	[0.80853359039673...	[0.72025491912707...	0.0
199810	18515	[0.24633418901301...	[-0.5591299711568...	1.0
199810	105760	[0.19777096337535...	[-0.7001422692047...	1.0

only showing top 6 rows

```
[84]: sort_gradient_boosted_tree_classifier_predictions_df =
↳ sort_gradient_boosted_tree_classifier_predictions.toPandas()
sort_gradient_boosted_tree_classifier_predictions_df.
↳ columns[sort_gradient_boosted_tree_classifier_predictions_df.isnull().any()]
```

```
[84]: Index([], dtype='object')
```

```
[85]: sort_gradient_boosted_tree_classifier_predictions_df.head()
```

```
[85]:   userID  trackID      probability \
0  199810   208019  [0.8085335903967381, 0.1914664096032619]
1  199810    74139  [0.8085335903967381, 0.1914664096032619]
2  199810     9903  [0.8085335903967381, 0.1914664096032619]
3  199810   242681  [0.8085335903967381, 0.1914664096032619]
4  199810    18515  [0.24633418901301743, 0.7536658109869826]

      rawPrediction  prediction
0  [0.7202549191270777, -0.7202549191270777]      0.0
1  [0.7202549191270777, -0.7202549191270777]      0.0
2  [0.7202549191270777, -0.7202549191270777]      0.0
3  [0.7202549191270777, -0.7202549191270777]      0.0
4  [-0.5591299711568518, 0.5591299711568518]      1.0
```

```
[86]: columns = ['userID', 'trackID']
sort_gradient_boosted_tree_classifier_predictions_df.
↳ to_csv('gradient_boosted_tree_classifier_predictions.csv', index=False,
↳ header=None, columns=columns)
```

```
[87]: gradient_boosted_tree_classifier_predictions_file =
↳ open('gradient_boosted_tree_classifier_predictions.csv')
final_gradient_boosted_tree_classifier_predictions_file =
↳ open('gradient_boosted_tree_classifier_submission.csv', 'w')
```

```
[88]: final_gradient_boosted_tree_classifier_predictions_file.
↳ write('TrackID,Predictor\n')
```

```
[88]: 18
```

```
[89]: last_user_id = -1
track_id_out_vec = [0] * 6
```

```
[90]: for line in gradient_boosted_tree_classifier_predictions_file:
    arr_out = line.strip().split(',')
    user_id_out = arr_out[0]
```

```

track_id_out = arr_out[1]

if user_id_out != last_user_id:
    i = 0

track_id_out_vec[i] = track_id_out

i = i + 1
last_user_id = user_id_out

if i == 6:
    predictions = np.ones(shape=(6))
    for index in range(0, 3):
        predictions[index] = 0

    for ii in range(0, 6):
        out_str = str(user_id_out) + '_' + str(track_id_out_vec[ii]) + ','
        ↪+ str(int(predictions[ii]))
        final_gradient_boosted_tree_classifier_predictions_file.
        ↪write(out_str + '\n')

```

```

[91]: gradient_boosted_tree_classifier_predictions_file.close()
      final_gradient_boosted_tree_classifier_predictions_file.close()

```